BLOG

AI Micro-Drama: How to Make Short Vertical Episodic Films with AI

By invideo

Last updated July 24, 2026

TL;DR

An AI micro-drama is a short vertical episodic story — typically 60–120 seconds per episode — produced end-to-end with generative video. Inside invideo, a creative producer agent breaks the script into scenes, a storyboard agent blocks each shot, and a DOP agent routes each shot to the right video model (Veo, Kling, Seedance 2.0), so a solo creator can ship a season without a crew.

> — produced end-to-end with generative video. Inside invideo, a creative producer agent holds the series bible, a storyboard agent blocks every shot, and a DOP agent routes each shot to the right video model (Veo, Kling, Seedance 2.0), so a two- or three-person team can ship a full season without a traditional crew.

What is an AI micro-drama?

A micro-drama is a vertical, phone-first drama series built from very short episodes — usually 1–3 minutes each — released in long seasons of 60–100 episodes, with every episode ending on a hook that pulls the viewer into the next one. An **AI micro-drama** is the same format produced with generative video instead of a physical shoot: the script is written or refined with AI, every shot is generated by a video model, and a crew of AI agents handles the direction layer — scene breakdown, storyboarding, shot routing, and continuity.

The format's grammar is fixed and worth internalizing before you produce anything: vertical framing in your delivery format, one dramatic beat per episode, a cliffhanger every 60–120 seconds, and a season long enough to sustain binge behavior. What changes with AI is not the grammar — it's who can execute it. A documented production shipped a 10-episode micro-drama series in 3 days with a 3-person crew, with each episode running approximately 1.5 to 2 minutes. That team size and timeline is the practical definition of what AI filmmaking now makes possible in this format.

Why AI changed micro-drama economics

A traditional micro-drama season costs $150k–$300k to shoot: cast, locations, crew days, and dozens of setups compressed into a punishing schedule — all before a single episode has proven it can retain viewers. The industry is projected to exceed $14B by 2027 per industry reporting, and the studios capturing that growth operate production: "For a Microdrama studio, the game isn't only a great script. It's shipping enough of them, fast enough, to find your winner as soon as possible." Generative video collapses the cost of each attempt. The same documented production measured its agent-driven pipeline at 5X faster than a standard micro-drama production pipeline, run by 3 people instead of a full crew. When a season no longer requires a six-figure shoot, you can test three concepts in the time and budget a traditional studio spends validating one — and kill the two that don't retain.

The economics also change what a "studio" is. The hierarchical agent structure described below compresses what would traditionally be a large production into a small-team operation, and the same structure scales the other direction: an agency managing multiple series for different clients runs one project per series, each with its own agent crew.

How to make Microdramas with AI Agents (Fil
Invideo for Filmmaking

8:25 / 10:16 Premiere Pro Plugin

More videos

End-to-end guide to making AI micro-dramas with agents

The end-to-end workflow: script to shipped episode

video models available, so the entire pipeline below runs in one place — the agents you create there are the crew. For a deeper pass on the scripting and scene-breakdown stage specifically, see the script to micro-drama workflow.

Step 1: Load the full script into a creative producer agent

Create one creative producer agent for the entire series and load it with everything: the show bible, tone, visual language, character descriptions, world rules, and the full season's scripts. This agent functions as the executive producer — as the documented production's lead put it: "I create one showrunner agent. Think of this as the executive producer of the entire series. It holds the big picture, the tone, the visual language, the character, world, everything." The operating principle is "Show Bible = single source of truth. Locked scripts = execution mode" — once scripts are locked, the invideo agent stops debating creative direction and starts executing it.

Before any video generation, run the cast and location build in parallel: generate character sheets for every character — front, side, and three-quarter angles plus facial close-ups — and reference sheets for every location, all in a single pass. Lock the finals and store them as shared context assets. Every later iteration — a costume change, a lighting shift, a new angle — starts from these locked base images, never from scratch. Skipping this lock before generating frames is the single most expensive ordering mistake in the format.

If a team is producing with you, run it as a shared-context operation: every team member loads the same locked show bible, episodic scripts, character sheets, and location references into their own agent instance. "Every team member uploads the same locked show bible and episodic scripts into their agent. No one's working of a different version."

Step 2: Break each episode into a shot list with a storyboard agent

Spin up one director-level sub-agent per episode, each with its own notebook page — the documented production flags this configuration step as critical: "This is important that you create a notebook with every episode director." The per-episode

Within each episode, run a storyboard agent to convert the script into a shot-by-shot breakdown. You don't build a shooting schedule manually — the invideo agent reads the episode script, understands its pacing and beats, and outputs the breakdown that becomes your production schedule: shot descriptions, blocking, and the order of generation. For high-stakes or complex scenes, storyboard first: lay out the key frames as a single vertical composite image (GPT-Image-2 handles vertical storyboard frames well), and attach that image to the video model for animation. This spends fewer credits and gives you more directorial control — "Your creative director can approve every storyboard frame before even a single credit is spent on video." The trade-off is real: a storyboard-first workflow sacrifices the unexpected shots a model generates freely, so reserve it for scenes where precision matters more than discovery. For insert shots and passing montage beats, skip location image generation entirely and prompt the model directly from the script description.

Step 3: Route each shot to the right video model

Assign a DOP agent to match each shot in the breakdown to the model that handles it best: Veo for cinematic realism and dialogue close-ups, Kling for multi-shot sequences it can generate natively, and Seedance 2.0 reference-to-video for any shot that must carry a character or a space forward from a previous shot. All of these models run inside the invideo AI video generator, so routing is a per-shot decision, not a platform decision.

Sequence generation with a spatial anchor: start every scene with a wide establishing shot, then generate each subsequent shot using the previous video as a spatial layout reference through Seedance 2.0 reference-to-video. "Every sequence starts with a wide establishing shot. The Agent holds the spatial layout. Every shot after references the last - same space, same lighting, same character position." This chaining is what keeps a generated scene feeling like one continuous space instead of a sequence of disconnected clips.

Generate 3–4 options per shot. That coverage is the documented best practice for giving your edit real choices — a single generation per shot leaves the editor cutting

Step 4: Assemble, extend, and cliffhang

Cut each episode from your selects, and use extend to stretch a strong 5-second generation into a full dramatic beat rather than regenerating the moment from scratch. Keep dialogue distributed across generations: cramming too much dialogue into a single 15-second generation breaks pacing and strips the editor of control over line timing. If you cut in a traditional NLE, the invideo plugin for Adobe Premiere Pro imports generated assets directly into a bin — a demonstrated workflow step pulled 8 assets in a single click, locally available the moment you import. A light upscale pass (Topaz Starlight 2.5 tested strongest on faces and fabrics in the documented production) and a film-grain pass in the NLE finish the footage.

Structurally, every episode ends on a hook — a reveal, a reversal, or an interrupted action that makes the next episode compulsory. Write the hook into the script stage, not the edit stage; the shot list should already know which shot is the cliffhanger. We break down hook mechanics episode-by-episode in the cliffhanger structure for micro-drama guide. Once episodes are shipping, the distribution question becomes commercial — here's how to monetize a micro-drama series once you have retention data.

Which AI model for which micro-drama shot?

Model choice is per-shot, not per-series, and the three roster video models divide the format's shot types cleanly:

Veo — cinematic realism: dialogue close-ups, emotional beats, shots where facial performance carries the episode.

Kling — multi-shot exchanges: it generates multi-shot sequences natively, which suits back-and-forth dialogue and action beats that need internal cuts.

Seedance 2.0 — series continuity: its reference-to-video mode takes a previous shot or a storyboard frame as a spatial and character reference, making it the holding the routing logic — you approve the shot list, it picks the model per shot. The full decision matrix, including when to override the default routing, is in which AI video model for which shot.

Keeping characters consistent across 60+ episodes

Character consistency across a long season comes from locked references, not from prompting skill. Lock the character sheets built in Step 1 — front, side, three-quarter, facial close-ups — before generating a single frame, and treat them as immutable base assets that every generation and every iteration starts from. Then let Seedance 2.0 on invideo carry those references into video: its reference-to-video mode holds face, wardrobe, and position against the locked sheets shot after shot.

The centrally locked context is what makes this survivable at 60+ episodes. Because every episode's director sub-agent pulls from the same central show bible and character sheets, a mid-series change propagates automatically: "Mid-series change like a scar or a haircut? Update the context once, the Agent remembers for every episode." You never re-brief sixty agents; you edit one source of truth.

When a single generation still misses — right face in seconds one to three, drift in the last two — assemble a Frankenstein shot: stitch the best seconds from multiple generations of the same shot into one clean take instead of regenerating until a full clip lands. The complete playbook, including reference-sheet specs and repair order, is in character consistency across episodes.

Vertical format: shooting for 9:16

Generate vertical from the first frame — prompt every shot in your delivery format rather than generating horizontal and cropping, because a crop discards the composition the model actually built. Vertical framing changes blocking: faces and upper bodies dominate the frame, two-character exchanges work as alternating platform UI.

The storyboard-first technique from Step 2 reinforces this — laying key frames out as a single vertical composite forces you to design in the delivery format before spending video credits. Shot-level prompting patterns for vertical are covered in generating vertical shots with AI, and if you're deciding runtime, the documented production settled on roughly 1.5–2 minutes per episode — the full reasoning is in how long should a micro-drama episode be.

What an AI micro-drama actually costs

The documented production's economics, quoted directly: "3 people. 3 days. 10 episodes. $1000 per episode." That is an all-in figure for an AI-assisted series with episodes running 1.5–2 minutes — roughly $500–$650 per finished minute for that team and approach. Costs vary with team, episode length, and how many options you generate per shot, so treat these as documented actuals, not a canonical number.

	Traditional micro-drama	AI micro-drama (documented production)
Season cost (60–100 episodes)	$150k–$300k	~$60k–$100k extrapolated from per-episode actuals
Per episode	~$1,500–$5,000 (season cost ÷ episode count)	~$1,000 all-in
Per finished minute	varies with shoot scale	~$500–$650 (1.5–2 min episodes)
Crew	full production crew	3 people
Timeline (10 episodes)	weeks of shoot + post	3 days

producing an AI micro-drama, and current credit plans are on the invideo pricing page.

Common mistakes when producing an AI micro-drama

Over-long episodes. The format's retention mechanics live at 60–120 seconds; a 4-minute episode is a short film wearing the wrong clothes. Cut the beat, not the runtime.

No cliffhanger. An episode that resolves cleanly ends the binge. Every script should identify its hook shot before generation starts.

Generating before locking references. Teams that lock character sheets after generating frames burn credits repairing drift. Lock first, generate second — and start every iteration from the locked base images.

One model for every shot. A dialogue close-up, a multi-shot exchange, and a continuity-critical follow-on shot are three different routing decisions (Veo, Kling, Seedance 2.0 respectively). Flat routing produces uneven episodes.

Treating each episode as a standalone project. Without a central show bible feeding per-episode sub-agents, continuity decays and your post team digs through hundreds of generations to find one shot. One series, one creative producer agent, one director sub-agent per episode.

Cramming dialogue into single generations. Too many lines in one 15-second clip breaks pacing and removes the editor's control over timing. Distribute dialogue across shots.

Iterating creative details mid-pipeline. Fixing creative decisions inside the generation loop expands production time because every fix is made along the way — lock scripts and references up front so the invideo agent runs in execution mode.

An AI micro-drama is a vertical, phone-first episodic drama — typically 60–120 seconds per episode across a 60–100 episode season — produced with generative video instead of a physical shoot. Script breakdown, storyboarding, model routing, and continuity are handled by a crew of AI agents; the shots themselves are generated by video models like Veo, Kling, and Seedance 2.0, all available inside invideo.

How long is a micro-drama episode?

Most episodes run 1–3 minutes, with 60–120 seconds the retention sweet spot. A documented AI production settled on approximately 1.5 to 2 minutes per episode across a 10-episode series. Each episode carries one dramatic beat and ends on a cliffhanger.

Which AI video model is best for micro-drama?

No single model covers the format — routing is per shot. Veo handles cinematic realism and dialogue close-ups, Kling generates multi-shot sequences natively, and Seedance 2.0 reference-to-video carries characters and spaces across consecutive shots. Inside invideo, a DOP agent makes that routing decision shot by shot.

How much does it cost to produce an AI micro-drama?

A documented production ran $1,000 per episode all-in — 10 episodes for roughly $10,000, produced by 3 people in 3 days — versus $150k–$300k for a traditionally shot season. Costs vary with episode length, option count per shot, and team, so treat published figures as documented actuals rather than a fixed rate.

Can one person produce a full AI micro-drama season?

Yes — the agent crew replaces most crew roles. The documented benchmark is a 3-person team shipping 10 episodes in 3 days at 5X the speed of a standard pipeline; a solo creator runs the same structure (one creative producer agent, per-

Lock character reference sheets — front, side, three-quarter angles plus close-ups — before generating anything, store them in central context every episode's sub-agent pulls from, and generate continuity shots with Seedance 2.0 reference-to-video. Mid-series changes are updated once in the central context and propagate to every subsequent episode; when one generation misses, stitch the best seconds from several generations into a single Frankenstein shot.

Sources

Micro-drama market projection ($14B+ by 2027): industry reporting and trade coverage of vertical short-drama platforms, widely aggregated in creator communities including Reddit.

Production figures (3-person crew, 3-day timeline, $1,000/episode, 5X pipeline speed, 1.5–2 minute episodes, 3–4 options per shot): invideo's documented micro-drama production.

SHARE

invideo AI

VIDEO TOOLS

AI Video Generator

AI Avatar Generator

Text to Video

Image to Video

UGC Ads

Image Tools

AI Image Generator

AI Art Generator

Image to Image

AI Thumbnail Maker

AI Image Enhancer

Google Veo 3.1

Sora 2

Kling AI

Wan AI

Pixverse AI

Hailuo AI

Seedance

View all

Flux Kontext

Qwen

Imagen AI

Reve AI

View all

Audio & Music Tools

AI Voice Generator

AI Music Generator

Text to Speech

Voice Cloning

AUDIO MODELS

Eleven Labs

Minimax

View all

QUICK AI TOOLS

Face Swap

Face Swap Video

invideo Studio

Online Video Editor

Picture Video Maker

Instagram Video Editor

Online Video Maker

Video Collage Maker

Video Trimmer

Video Compressor

MP4 Compressor

Help

Email Us

Tutorials

Blogs

Youtube Channel

Instagram Inspiration

Community

GET THE APP

GET IT ON Google Play

Download on the App Store