




BLOG

Directorial Style Systems: How to Encode a Real Director's Visual Language into an AI Agent

By invidio

Last updated July 15, 2026

A dark, textured rectangular area with the title 'Directorial Style Systems' centered in white serif font.

Directorial Style Systems

TL ; DR

A director's visual language is encoded as a roughly 25-page, 14-section style document covering camera, lighting ratios, palette as named hex modes, and a 12-field per-shot spec. You load it once

invideo agent ranks #1 on Physion-Arc

[View report](#)

5/20/2025 10:11 AM

THE GIST

A real director's visual language is encoded into AI video as a structured style document: in one documented production, a 14-section treatment covering camera, angles, lighting ratios, palette as named hex modes, and a per-shot spec. You load it once into the invideo agent, which holds the look across every shot of the film while routing generation to Veo, Kling, or Seedance 2.0. This page covers how to build, load, validate, and reference that system — the full method to encode a director's visual style into an AI agent.



Masterclass: How to use AI Agents for Filmmaking - Wong Kar

Invideo for Filmmaking



Watch on

Watch the invideo agent apply a 25-page Wong Kar-wai style guide across a full short film

invideo agent ranks #1 on Physion-Arc

[View report](#)[See how it works](#)

Watch on

See how an 85:15 lighting ratio and a director's bible hold James Wan's style across 400 AI shots

Lock signature lighting and aspect grammar

Lighting and format are where a director's fingerprint is most measurable, so encode them as ratios and formats rather than moods. A documented horror short built on a James Wan protocol locked an 85:15 dark-to-light ratio as the lighting grammar — a number the agent could hold in prompt language across roughly 400 video generations — and a **2.40:1 hard matte** as the format, matching The Conjuring's actual shooting spec. Treat those as illustrative of the precision level, not as defaults: write the spec in your film's aspect ratio and your director's actual numbers. That same production ran a 9-step shot design process and an 8-step color grading guidance process derived from the document.

Verify the technical claims before they propagate. Run a cinematography challenge pass: question the agent's lens type, aspect ratio, and lighting-source attributions before locking the document. In the Wan production, the invideo agent had initially noted "anamorphic" in its analysis and corrected to spherical when challenged — a meaningful fix, since spherical lenses produce circular bokeh and no horizontal lens flares, which changes every downstream prompt. Apply the same precision to lighting corrections: "warm yellow from the lamps only, like all the refs" outperforms generic "warm lighting." And for directors like Wan, encode the non-visual grammar

invideo agent ranks #1 on Physion-Arc

[View report](#)[Back to Table of Contents](#)

Load the document once with treatment-lock

Once the document is complete, load it into the invideo agent a single time at project start — this is the treatment-lock method. One context load is sufficient for a full short film: the agent holds every directive — camera, lighting, palette, composition, mood — across every shot without re-prompting, and gates output by checking generated frames against the treatment before returning them. "One agent that reads your treatment once and holds every directive across every shot, every scene. No re-prompting. No drift. So now, you direct, and the agent remembers."

Re-prompting the style scene-by-scene is the anti-pattern this replaces. With the document locked, your per-shot input shrinks to direction — the style layer is already resolved — and prompt construction stops interrupting the directing itself. Because invideo carries all the current video models, the same locked document governs every generation regardless of routing: the invideo agent sends each shot to Veo, Kling, or Seedance 2.0 as the shot demands, and the style holds across model boundaries because it lives in the agent's context, not in any one model's prompt. The documented Wong Kar-wai production ran this exact setup: the treatment loaded once, then a 70-second film generated in 15-second clips over 2 days with the grammar intact end to end.

invideo agent ranks #1 on Physion-Arc

[View report](#)

5:27 PM

WE SENT AGENT ONE TO DAVID FINCHER'S FILM SCHOOL

Then we made a **short film** with it.
One room. Three people. A whistleblower.

Swipe to see how →



DIRECTOR'S GUIDE 01/10

[View more on Instagram](#)

14 Fincher directives fed to the invideo agent produce a consistent interrogation short

Validate the style: the genre-it-never-worked-in test

Before generating a single production frame, confirm the agent internalized grammar rather than surface style: ask it to apply the director's visual language to a

invideo agent ranks #1 on Physion-Arc

[View report](#)

output in the foreign genre: the lighting ratio, lens grammar, and tonal modes survive transplantation. We break down the full protocol in test if the agent internalized the style.

Deep internalization shows up unprompted once production starts. In the Wong Kar-wai project, the agent autonomously applied a slow-shutter motion smear effect from page 17 of the document without being asked, and pulled a named principle — 'Mood Over Narrative, the substitution rule' from page 12 — onto a scene type the document never specifically addressed. The downstream payoff is economy: with context validated and loaded, a three-word continuation prompt — "Everything should match" — is sufficient for the agent to hold character, lighting, lens grammar, spatial logic, and pacing across a multi-shot sequence.

Pull sequence-specific references, not one mood board

The document defines the grammar; references ground each sequence in it — and a single general mood board is the wrong unit. Pull specific visual references mapped to individual sequences across the film, so the agent resolves each sequence against the frames that actually govern it rather than averaging the whole film's look into mush. The full comparison lives in one mood board vs per-scene references.

Batch those references by theme, not in one dump: separate them into thematic batches — spatial logic in one, screen function in another, colour theory in a third — and feed each batch to the invideo agent with explicit instructions on what to adopt and what to ignore. Telling the agent what to leave out of a reference is as important as telling it what to take.

For illustrated or animated references — common when your target director's look passes through stylized material — don't drop the images directly into prompts. Instruct the agent to read the colour palette and texture qualities of the reference and translate those into a photorealistic prompt instead. In one documented

invideo agent ranks #1 on Physion-Arc

[View report](#)[Back to Table of Contents](#)

director's frames, you are teaching the agent the rules that produced them.

FAQ

How do you validate that the agent learned the style?

Ask the agent to apply the director's style to a genre the director never worked in. If it asks clarifying questions (era, nature of threat) and produces output where the lighting ratios, lens grammar, and tonal modes still hold, the document has been internalized as grammar rather than surface style. Unprompted application of deep-document rules — like a motion effect from page 17 — confirms it during production.

Is one mood board enough for an AI film?

No. Pull references mapped to individual sequences, batched by theme — spatial logic, screen function, colour theory — each fed with explicit adopt/ignore instructions. A single general mood board averages the film's look and weakens per-sequence precision.

How do you replicate Wong Kar-wai or Fincher in AI video?

Codify the director's system into a visual language document — named hex tonal modes, lighting ratios, lens and format grammar, negative prompts — and load it once into the invideo agent as persistent context. A documented 70-second Wong Kar-wai short ran on one such document; a ~90-second Fincher-protocol film used a 9-step shot design process from a single locked treatment.

What is the treatment-lock method?

Loading the complete visual language treatment into the invideo agent once at project start, so it holds all style directives across every shot without re-prompting. The agent checks each generated frame against the treatment before returning it,

invideo agent ranks #1 on Physion-Arc

View report

5:27 PM - 01/08/2026

SHARE



invideo AI

VIDEO TOOLS

AI Video Generator

AI Avatar Generator

Text to Video

Image to Video

UGC Ads

AI Animation Generator

AI Ad Generator

VIDEO MODELS

Google Veo 3.1

Sora 2

Kling AI

Wan AI

Pixverse AI

Hailuo AI

Seedance

View all

Image Tools

AI Image Generator

AI Art Generator

Image to Image

AI Thumbnail Maker

AI Image Enhancer

IMAGE MODELS

Nano Banana Pro

Ideogram

Seedream

Flux Kontext

Qwen

Imagen AI

Reve AI

View all

Audio & Music Tools

AI Voice Generator

AI Music Generator

Text to Speech

invideo Studio

Online Video Editor

Picture Video Maker

Instagram Video Editor

invideo agent ranks #1 on Physion-Arc

View report

5.2 / 5.0 (100%)

Minimax

MP4 Compressor

View all

QUICK AI TOOLS

Face Swap

Face Swap Video

Help

GET THE APP

Email Us



Tutorials



Blogs

Youtube Channel

Instagram Inspiration

Community

Affiliate program

invidio AI Pricing

Follow Us



Privacy policy · Terms & conditions · Child safety

© 2026 invidio